

# Glibc Heap Allocator

CS390R - UMass Amherst

# Course Information

- Homework Released today

# Today's Content

- Glibc Heap
- Safely using the Heap
- Heap Chunks
- Free lists
- Allocation & Freeing procedures

# Glibc

- “GNU C Library”
- Core libraries for many systems using the Linux kernel
- Standard functions and wrappers around syscalls
- Examples:
  - strlen
  - printf
  - read/write
  - malloc
  - free
  - ...

# Malloc

- Glibc Memory Allocator
  - Exploited since 20 years - new techniques keep coming out
  - Provides process with dynamic memory
- 
- `malloc()` - allocate some memory
  - `free()` - free some memory (give it back)
  - `realloc()` - resize a previous allocation
  - `calloc()` - allocate some memory and zero it out

Lets walk through an  
Example

# PWNDbg Heap Commands


- vis - visualize heap chunks
- bins - print out various bins
- fastbins - print out just fast bins
- tcachebins - print out just tcache bins
- smallbins - print out just small bins
- largebins - print out just large bins
- unsortedbin - print out unsorted bin
- vmmap - print out memory mapping
- top\_chunk - print out address/size of top chunk

# Basic Rules when using the heap

- ❖ Do not read/write to a pointer after its been free'd
  - Use after free vulnerability
- ❖ Do not use uninitialized heap space
  - Information leaks
- ❖ Do not read/write after the max size of an allocation
  - Heap Overflow
- ❖ Do not read/write before the beginning of an allocation
  - Heap Underflow
- ❖ Do not free a pointer more than once
  - Double Free
- ❖ Do not free a pointer that did not originate from malloc
  - Invalid Free

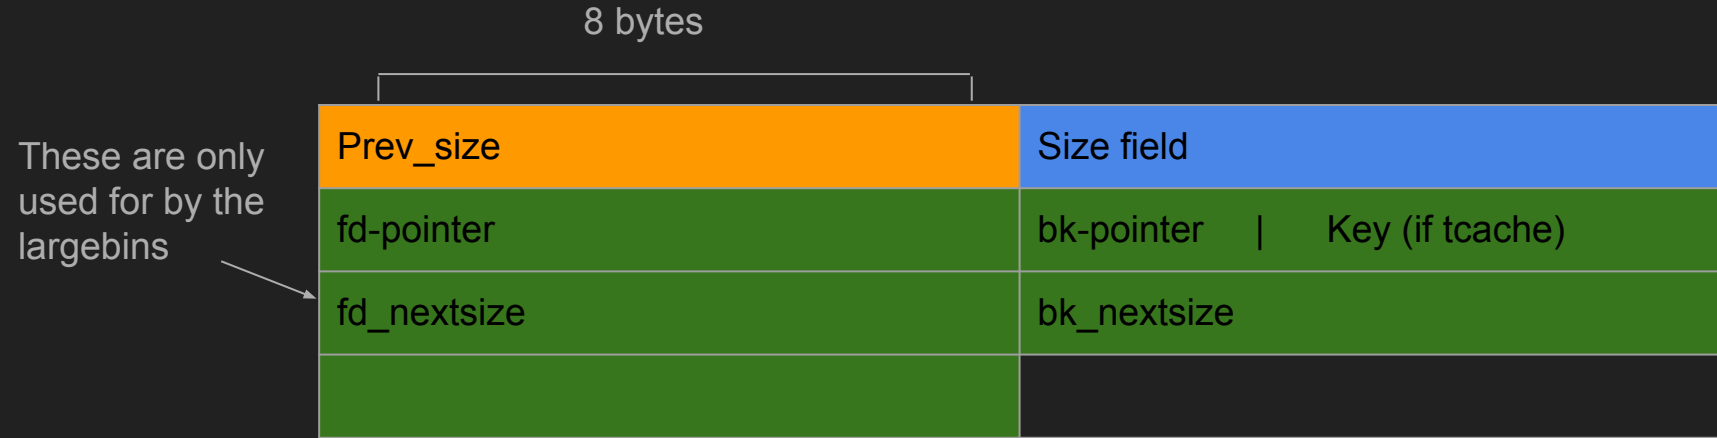


# Allocated Heap Chunks



- Minimum chunk size is 32
- Chunk sizes increment by 16 bytes
- Allocations are always 16-byte aligned
- Chunks are in 1 of 2 states, in-use or free

# Free'd Heap Chunks



- Some of the userdata-space is repurposed as metadata
- Metadata is used for linked lists
  - fd -> used by all types of free chunks
  - bk -> used by some types eg. smallbins (tcache uses this place for a key)
  - nextsize-fields -> only used by largebins
  - Prev\_size field only used in chunks that support consolidation

# Arenas

- For multithreaded applications, arena's are used to give each thread a separate heap
- This lets different threads use their own heaps without needing locks
- These arena's cannot use `sbrk()` to grow their heaps, so they make use of `mmap()` + `mprotect()` to emulate this behavior
- Chunks that are not in the main arena have the appropriate flag set in their metadata

# Free lists

- Heap manager tracks free chunks in different free lists called “bins”
- Allocation request -> retrieve a chunk from one of these bins if possible
- Multiple different lists are used for performance reasons:
  - Tcache
    - 64 tcache bins with sizes [0x20-0x410]
  - Fastbins
    - 7 fastbins [0x20 - 0x80]
  - Unsorted bin
    - 1 unsorted bin
  - Small bins
    - 62 small bins with sizes [0x20 - 0x3f0]
  - Large bins
    - 63 large bins with sizes [0x400 - 0x80000]


# Simplified Chunk Allocation Strategy

- If there is a big enough chunk in bins, use it for the new allocation
- No free chunks -> allocate memory from top chunk
  - Decreases the size value & place the new chunk at the end of the current heap region
- Top chunk does not have enough space -> kernel required
  - Sbrk() system call to ask kernel to grow the heap region
- If heap has reached the max possible size, mmap() -> create an entirely new heap somewhere else in memory
- If everything fails, return NULL
  
- Exception for very large allocations, immediately mmap() it. When they are later free'd, they are completely removed instead of being placed into any bins

# Simplified Freeing Strategy

- If the IS\_MMAPED flag is set, immediately unmap it
- If in tcache range, link into appropriate tcache bin
- If in fastbin range, link into appropriate fastbin
- Attempt to consolidate (merge with other free chunks next to it)
- Link into unsorted bin

# Arena




# Fast Bins - 1

- Fast bins are the second bins accessed after the tcache bins during allocation
- Keep recently free'd chunks on a fast turnaround queue via singly linked list
- Each fastbin is only responsible for a single sized chunk (0x20, 0x30, ... 0x80)
- Since everything is sorted, insertions and removals are very fast
- Chunks linked into the fastbins are not consolidated with other chunks through standard means (There are some exceptions)
- The head of each fastbin resides in its arena




# Fast Bins - 2

- Linked list pointers are stored as inline metadata within the chunk
- First 8 bytes of a chunk's userdata are used as a forward pointer
- LIFO structure
- Unlinking a chunk from singly linked fastbin
  - copy victim chunks fd into the head of the list



# Small Bins

- 62 of them with sizes [0x20 - 0x3f0]
- One size/bin for fast insertions/deletions
- Head of each small bin in arena but inline links
- Small bins are searched after tcache, fastbin, and unsorted bins
  - if the request falls in small bin size range
- Doubly linked circular list - FIFO




# Large Bins

- 63 of them with sizes [0x400 - 0x80000]
- Store chunks within size ranges
- Insertions require traversing the list -> slower than other bins
- Head of each small bin in arena but inline links
- Large bins are searched after tcache, fastbin, and unsorted bins
  - if the request falls in large bin size range
- Doubly linked circular list - Since the list has potentially multiple copies of various sizes, and the list has to be traversed, a skip list is used so multiple chunks of one size aren't all traversed



# Unsorted Bin

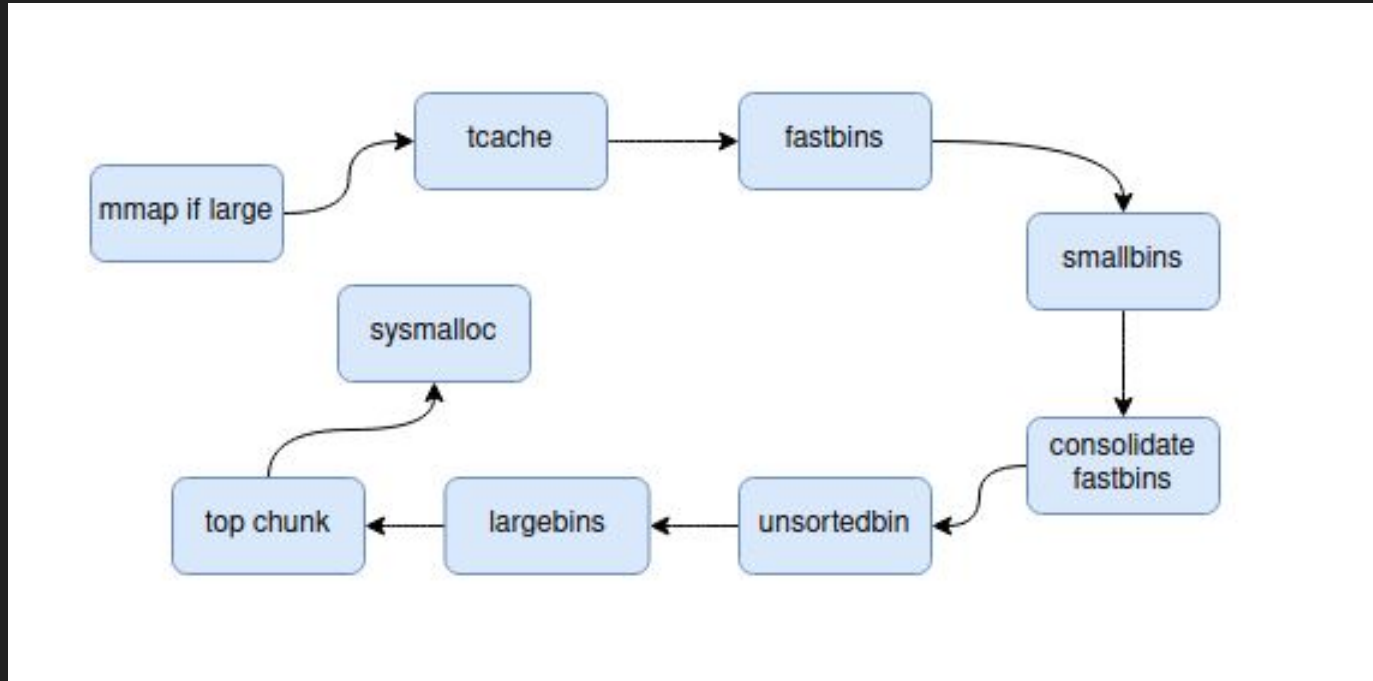
- Free's that would go into small/large bins go into unsorted bin first
- Allocations traverse unsorted bin looking for fitting chunk, before looking in small/large bins depending on size
- During this traversal, all chunks are sorted into small or large bins
- Circular doubly linked list



# Tcache (thread-cache)

- Relatively new addition to glibc (Glibc versions  $\geq 2.26$ )
- Limit on concurrent arenas, but each thread gets its own tcache
  - Reduces time spent waiting on other threads
- 64 of them with sizes [0x20 - 0x410]
- Holds up to 7 chunks per size field
- During allocations, the tcache is the first set of bins to be checked
- Similarly to fast bins, not regularly coalesced - LIFO
- Has a “count” bitmap-like structure with a field for each tcache size, and a pointer within itself that points to first entry of each bin

# Full Allocation Strategy



# Full Freeing Strategy

