

Glibc Heap Exploitation

CS390R - UMass Amherst

Course Information

- Project 4 still ongoing
- Homework due tonight

Today's Content

- Browsing Malloc Source
- Heap Exploits basics
- Hooks
- System vs One Gadget
- House of Force
- Tcache Dup

Browsing glibc source

- <https://www.gnu.org/software/libc/>
- <https://code.woboq.org/userspace/glibc/>

- malloc.c
 - :3035 void *__libc_malloc(size_t bytes);
 - :3093 void __libc_free(void *mem);

Heap Exploits

- Stack Exploits -> ROP
- Heap Exploits -> ???
- Write-What-Where primitives
 - Write an arbitrary value to an arbitrary location (similar to format string writes)
- Where to write
 - Fini-Array -> Partial Relro makes it non-writable
 - GOT -> Full Relro makes it non-writable
 - Hooks -> Possible until very recent glibc 2.34
- What to write
 - system()
 - One Gadget

Hooks /1

- malloc, free, realloc, etc all have hooks in memory
- When malloc is executed it first checks if there is a listed hook. If there is, the hook is executed instead.
- Debugging feature

```
void * __libc_malloc (size_t bytes) {  
    mstate ar_ptr;  
    void *victim;  
  
    _Static_assert (PTRDIFF_MAX <= SIZE_MAX / 2, "PTRDIFF_MAX is not more than half of SIZE_MAX");  
  
    // Read the malloc hook and execute it if found  
    void *(*hook) (size_t, const void *) = atomic_forced_read (__malloc_hook);  
    if (__builtin_expect (hook != NULL, 0))  
        return (*hook)(bytes, RETURN_ADDRESS (0));  
}
```

Hooks /2

```
pwndbg> p &_malloc_hook
$5 = (void *(*)(size_t, const void *)) 0x7ffff7f9bb70 <__malloc_hook>
pwndbg> dq 0x7ffff7f9bb70
00007ffff7f9bb70  00007ffff7e4cc90 0000000000000000
00007ffff7f9bb80  0000000000000000 0000000000000000
00007ffff7f9bb90  0000000000000000 0000000000000000
00007ffff7f9bba0  0000000000000000 0000000000000000
pwndbg> x 0x7ffff7e4cc90
0x7ffff7e4cc90 <malloc_hook_ini>: 0xfa1e0ff3
```

- We can print out the location/content of these hooks in gdb
- What if we overwrite one of these hooks using our exploit?

System()

- Goal is to call `system("/bin/sh");`
- Arg1 of function (eg. `malloc/free`) also becomes arg1 of corresponding hook
- Via `__malloc_hook`
 - `malloc(<address of "/bin/sh" string>) -> system("/bin/sh");`
- Via `__free_hook`
 - `free(<chunk that has "/bin/sh" in its userdata>) -> system("/bin/sh");`

One Gadget

- `execve("/bin/sh")` call located in `libc`
- https://github.com/david942j/one_gadget
- Usually has some constraints that need to be met to successfully spawn a shell

```
→ glibc_heap_exploitation_1 git:(main) X one_gadget libc.so.6
0x4f3d5 execve("/bin/sh", rsp+0x40, environ)
constraints:
  rsp & 0xf == 0
  rcx == NULL

0x4f432 execve("/bin/sh", rsp+0x40, environ)
constraints:
  [rsp+0x40] == NULL

0x10a41c execve("/bin/sh", rsp+0x70, environ)
constraints:
  [rsp+0x70] == NULL
```

Patching in custom libc

- Project 5, the challenges will ship with custom libc versions
 - The binaries will already be patched with them since edlab does not have some of the required tools
-
- `patchelf --set-interpreter ./ld-2.31.so ./chal`
 - `patchelf --set-rpath ./ ./chal`

House Techniques

- <http://phrack.org/issues/66/10.html>
- Published in 2005
- Theoretical representation of new skills representing “modern” heap exploitation

4 - DES-Maleficarum...

4.1 - The House of Mind

4.1.1 - FastBin Method

4.1.2 - av->top Nightmare

4.2 - The House of Prime

4.2.1 - unsorted_chunks()

4.3 - The House of Spirit

4.4 - The House of Force

4.4.1 - Mistakes

4.5 - The House of Lore

4.6 - The House of Underground

House of Force /1

- Goal: Return arbitrary pointer from malloc
- Exploits top chunk size field corruptions
- Very large allocation
 - Usually this would get mmaped to a random location since its too large for the main arena



A screenshot of a memory dump with a black background and red and green text. The dump shows two rows of memory addresses and permissions. The first row shows a green address 0x405000, a green address 0x426000 with permissions 'rw-p', a green value 21000, and a green label '[heap]'. The second row shows a red address 0x7ffff7426000, a red address 0x7ffff7db0000 with permissions 'rw-p', a red value 98a000, and a red label '[anon_7ffff7426]'. Below the first row, there is a faint, partially visible line of text that appears to be '0x7ffff7426000'.

0x405000	0x426000 rw-p	21000	[heap]
0x7ffff7426000	0x7ffff7db0000 rw-p	98a000	[anon_7ffff7426]

- What if the set the top chunk size to a very large value?
 - Suddenly we can make an arbitrarily sized allocation starting at the top chunk
- With this we can move the top chunk to an arbitrary location in memory.
- This means that we can control where the next allocation is made

Memory Map

```
pwndbg> vmmap
```

```
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
```

0x3ff000	0x400000	rw-p	1000	0	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x400000	0x401000	r--p	1000	1000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x401000	0x402000	r-xp	1000	2000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x402000	0x403000	r--p	1000	3000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x403000	0x404000	r--p	1000	3000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x404000	0x405000	rw-p	1000	4000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/demo_bin
0x405000	0x426000	rw-p	21000	0	[heap]
0x7ffff79e2000	0x7ffff7bc9000	r-xp	1e7000	0	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/libc.so.6
0x7ffff7bc9000	0x7ffff7dc9000	---p	200000	1e7000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/libc.so.6
0x7ffff7dc9000	0x7ffff7dcd000	r--p	4000	1e7000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/libc.so.6
0x7ffff7dcd000	0x7ffff7dcf000	rw-p	2000	1eb000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/libc.so.6
0x7ffff7dcf000	0x7ffff7dd3000	rw-p	4000	0	[anon_7ffff7dcf]
0x7ffff7dd3000	0x7ffff7dfc000	r-xp	29000	0	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/ld-2.27.so
0x7ffff7dfc000	0x7ffff7ffa000	rw-p	2000	0	[anon_7ffff7ff4]
0x7ffff7ffa000	0x7ffff7ffa000	r--p	4000	0	[vvar]
0x7ffff7ffa000	0x7ffff7ffc000	r-xp	2000	0	[vdso]
0x7ffff7ffc000	0x7ffff7ffd000	r--p	1000	29000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/ld-2.27.so
0x7ffff7ffd000	0x7ffff7ffe000	rw-p	1000	2a000	/home/seal/github/re-vr-course/lectures/glibc_heap_exploitation_1/ld-2.27.so
0x7ffff7ffe000	0x7ffff7fff000	rw-p	1000	0	[anon_7ffff7ffe]
0x7ffff7fff000	0x7ffff7fff000	rw-p	21000	0	[stack]
0xffffffffff600000	0xffffffffff601000	--xp	1000	0	[vsyscall]

```
pwndbg>
```

House of Force /2

- Technique is fully mitigated in Glibc 2.29 so we will be using glibc 2.27

Steps:

1. Overwrite topchunk with -1 (0xffffffffffffff)
2. Calculate size of our allocation:
 - a. If we need to wrap around VA: $0xffffffffffffff - TOP + Target - 0x20$
 - b. if we don't need to wrap around VA: $Target - 0x20 - TOP$
3. Allocate a chunk that overlaps the target

Tcache Dup

- Goal: Allocate chunk at arbitrary address
- Multiple ways to exploit, a popular one is a double free (glibc <= 2.29)
 - Free the same chunk twice to link it into tcache bins twice
 - Next 2 allocations will return the same chunk so you can overwrite metadata of the 2nd
 - This can be used to link a fake chunk into the tcache bins
 - The third allocation will then be made at the specified location

-> Write-What-Where vulnerability

```
0x405a90      0x0000000000000000      0x0000000000000021      .....!.....
0x405aa0      0x4141414141414141      0x0000000000000000      AAAAAAAA..... <-- tcachebins[0x20][0/2]
0x405ab0      0x0000000000000000      0x00000000000020551      .....Q..... <-- Top chunk

pwndbg> bins
tcachebins
0x20 [ 2]: 0x405aa0 ← 'AAAAAAA'
```